

# McMaster University



SEP 721 – Data Analytics, Machine Learning and AI on  
Cloud Platforms

Assignment 3: Azure Machine Learning

Submitted by,

Greeshma Gopal(gopalg)

ID- 400245291

## 1. Creating Azure Machine learning

- Creating work space in azure cloud for machine learning

Main \* Tags Review \*

Workspace Name \*

 ✓

Subscription

 ▼

Resource group

 ▼

[Create new](#)

Location

 ▼

Workspace edition [View full pricing details](#) ⓘ

 ▼

ⓘ For your convenience, these resources are added automatically to the workspace, if regionally available: [Azure storage](#), [Azure Application Insights](#) and [Azure Key Vault](#).

Review + Create

 Microsoft.MachineLearningServices | Overview

Deployment

Search (Cmd +/)

Delete Cancel Redeploy Refresh

Overview Inputs Outputs Template

✓ Your deployment is complete

Deployment name: Microsoft.MachineLearningServices Start time: 4/6/2020, 5:19:08 PM  
Subscription: Free Trial Correlation ID: 19ca63a5-aec2-411f-a7f5-f9530120afft  
Resource group: Greeshma\_resource

Deployment details [\(Download\)](#)

Next steps

Go to resource

- Creating compute instance in Azure machine learning studio

## New Compute Instance

Customers should not include personal data or other sensitive information in fields marked with because the content in these fields may be logged and shared across Microsoft systems to facilitate operations and troubleshooting. [Learn more](#)

Compute name \*

Region \*

Virtual Machine size \*

Enable SSH access

Advanced settings

Show created by me only Search to filter items...

Name	Status	Application URI	Virtual Machine size	Created on ↓
machinelearninginazure	Running	<a href="#">JupyterLab</a> <a href="#">Jupyter</a> <a href="#">RStudio</a> <a href="#">SSH</a>	STANDARD_D3_V2	Apr 6, 2020 6:53 PM

- Ensuring that the latest version of Azure machine learning is installed

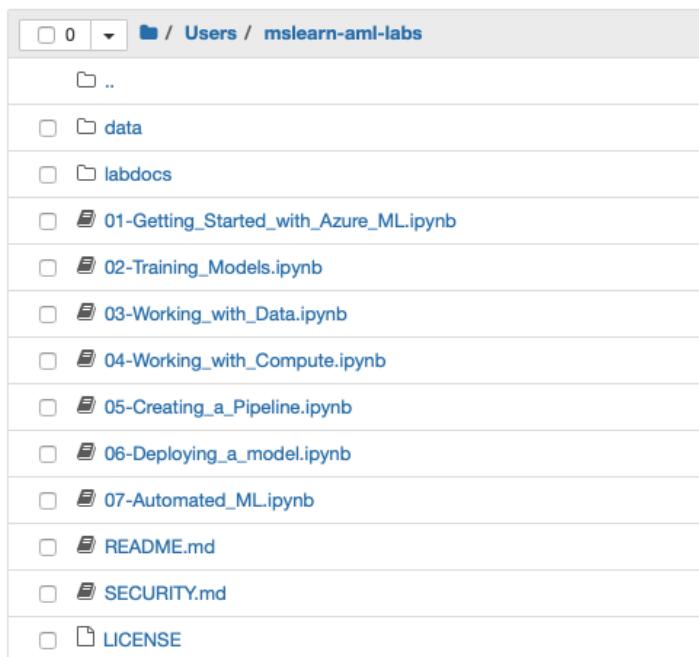
```
Successfully uninstalled azureml-interpret-1.0.85
Found existing installation: azureml-explain-model 1.0.85
Uninstalling azureml-explain-model-1.0.85:
Successfully uninstalled azureml-explain-model-1.0.85
Found existing installation: azureml-train-automl-runtime 1.0.85.5
Uninstalling azureml-train-automl-runtime-1.0.85.5:
Successfully uninstalled azureml-train-automl-runtime-1.0.85.5
Found existing installation: azureml-train-automl 1.0.85
Uninstalling azureml-train-automl-1.0.85:
Successfully uninstalled azureml-train-automl-1.0.85
Found existing installation: nbconvert 5.4.1
Uninstalling nbconvert-5.4.1:
Successfully uninstalled nbconvert-5.4.1
Found existing installation: papermill 2.0.0
Uninstalling papermill-2.0.0:
Successfully uninstalled papermill-2.0.0
Found existing installation: azureml-contrib-notebook 1.0.85
Uninstalling azureml-contrib-notebook-1.0.85:
Successfully uninstalled azureml-contrib-notebook-1.0.85
Found existing installation: azureml-widgets 1.0.85.1
Uninstalling azureml-widgets-1.0.85.1:
Successfully uninstalled azureml-widgets-1.0.85.1
Found existing installation: azureml-sdk 1.0.85
Uninstalling azureml-sdk-1.0.85:
Successfully uninstalled azureml-sdk-1.0.85
Successfully installed azure-core-1.4.0 azure-identity-1.3.1 azureml-automl-core-1.2.0 azureml-automl-runtime-1.2.0 azureml-contrib-notebook-1.2.0 azureml-core-1.2.0.post2 azureml-dataprep-1.3.6 azureml-dataprep-native-14.1.0 azureml-defaults-1.2.0 azureml-explain-model-1.2.0 azureml-interpret-1.2.0 azureml-pipeline-1.2.0 azureml-pipeline-core-1.2.0 azureml-pipeline-steps-1.2.0 azureml-sdk-1.2.0 azureml-telemetry-1.2.0 azureml-train-1.2.0 azureml-train-automl-1.2.0 azureml-train-client-1.2.0 azureml-train-automl-runtime-1.2.0 azureml-train-core-1.2.0 azureml-train-restclients-hyperdrive-1.2.0 azureml-widgets-1.2.0 flask-1.0.3 future-0.18.2 interpret-community-0.7.0 interpret-core-0.1.20 msal-1.2.0 msal-extensions-0.1.3 nbconvert-5.6.1 papermill-1.2.1 portalocker-1.6.0 scipy-1.1.0 werkzeug-0.16.1 wheel-0.30.0
azureuser@azurermachinelearning:/mnt/batch/tasks/shared/LS_root/mounts/clusters/azurermachinelearning/code$
```

- Cloning the github code

```
https://github.com/microsoftdocs/mslearn-aml-labs
Cloning into 'mslearn-aml-labs'...
remote: Enumerating objects: 147, done.
remote: Counting objects: 100% (147/147), done.
remote: Compressing objects: 100% (103/103), done.
remote: Total 147 (delta 89), reused 83 (delta 43), pack-reused 0
Receiving objects: 100% (147/147), 412.51 KiB | 0 bytes/s, done.
Resolving deltas: 100% (89/89), done.
Checking connectivity... done.
```

- The files have been cloned in the folder ‘Users’

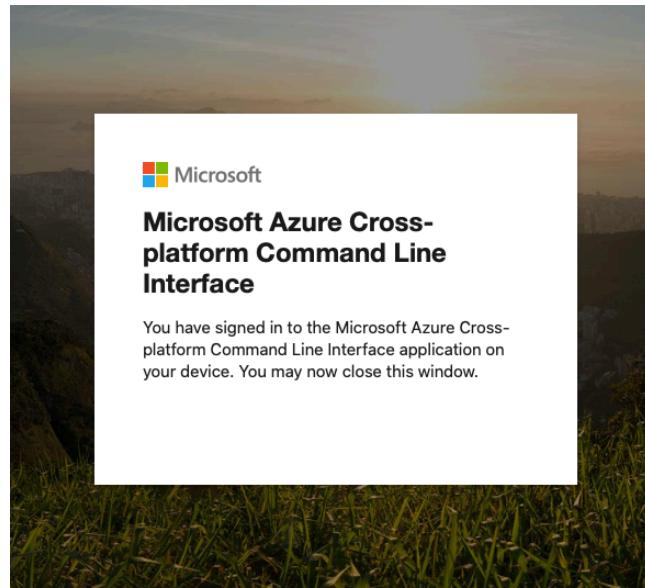
Select items to perform actions on them.



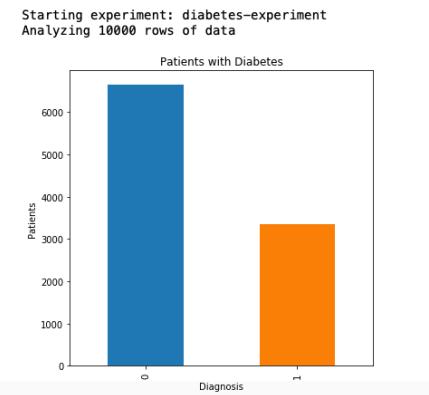
- Importing libraries and logging in the command line interface

```
[2]: from azureml.core import Workspace
ws = Workspace.from_config()
print(ws.name, "loaded")
```

Performing interactive authentication. Please follow the instructions  
 To sign in, use a web browser to open the page <https://microsoft.com/d>  
 authenticate.  
 Interactive authentication successfully completed.  
 greeshma1 loaded



- Run an experiment to explore the data, extracting statistics, visualizations, and data samples

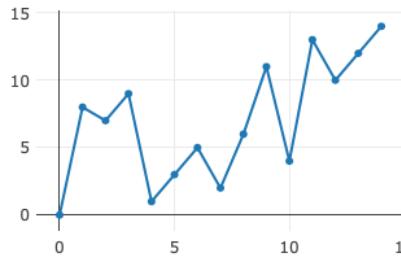


- Run object to get information about the run and its outputs

```
files = run.get_file_names()
print(json.dumps(files, indent=2))
{'aml://gitlab.com/microsoft/labs/mllearn-labs', 'mlflow.source.git.repoUrl': 'https://github.com/microsoft/labs/mllearn-aml-labs', 'azureml.git.branch': 'master', 'mlflow.source.git.branch': 'master', 'azureml.git.commit': 'f33990a44ca85bedb6961c50368c92d0f7', 'mlflow.source.git.commit': 'f33990a4fa96e44ca85bedb6961c50368c92d0f7', 'azur
eml.git.dirty': 'False', 'ContentSnapshotId': '96f1383d-29ac-488c-91c1-5b28bb120c30'}, 'inputDatasets': [], 'logFil
es': {}}
{
  "observations": 10000,
  "label distribution": "aml://artifactId/ExperimentRun/dcida.69c0b7ca-4c21-4430-9635-c2b8ec9412b1/label distributio
n_1586214726.png",
  "pregnancy categories": [
    0,
    8,
    7,
    9,
    1,
    3,
    5,
    2,
    6,
    11,
```

Run is completed.

pregnancy categories



greeshma1 > Experiments > diabetes-experiment > Run 1

Run 1 ✓ Completed

↻ Refresh ↻ Resubmit ✖ Cancel

Details Metrics Images Child runs Outputs + logs Snapshot Raw JSON Explanations (preview)

#### Properties

Status  
Completed  
Created  
Apr 6, 2020 7:11 PM  
Duration  
15.29s  
Compute target  
sdk  
Run ID  
69c0b7ca-4c21-4430-9635-c2b8ec9412b1  
Run number  
1  
Script name

#### Metrics

label distribution  
Type: Image  
BMI  
Type: Table  
SerumInsulin  
Type: Table  
TricepsThickness  
Type: Table  
DiastolicBloodPressure  
Type: Table  
PlasmaGlucose  
Type: Table  
pregnancy categories

- Creating folder for the experiment files, and copy the data into it

```
os.makedirs('outputs', exist_ok=True)
data.sample(100).to_csv("outputs/sample.csv", index=False, header=True)
```

```
# Complete the run
run.complete()
```

```
Writing diabetes-experiment-files/diabetes_experiment.py
```

- Creating conda environment to run the experiment 2

33\_2e4724ff

Arguments	N/A
Label:0	6656
Label:1	3344
observations	10000

**Output Logs** click to scroll output; double click to hide [logs/azureml/8590\\_azureml.log](#)   Auto-switch

```
2020-04-06
23:19:00,390|azureml._history.utils.context_managers.FileWatcher.UploadQueue|DEBUG|Finished
calling finish
2020-04-06
23:19:00,390|azureml._history.utils.context_managers.FileWatcher|DEBUG|FileWatcher task queue
finished, exiting
2020-04-06 23:19:00,391|azureml._history.utils.context_managers.FileWatcher.10d8c378-d009-
4b29-b445-56dc4126c6e3|DEBUG|Using basic handler - no exception handling

Run is completed.
```

greeshma1 > Experiments > diabetes-experiment > Run 2

Run 2 ✓ Completed

↻ Refresh ↻ Resubmit ✖ Cancel

Details	Metrics	Images	Child runs	Outputs + logs	Snapshot	Raw JSON	Explanations (preview)
<b>Properties</b>				<b>Metrics</b>			
<p>Status Completed</p> <p>Created Apr 6, 2020 7:18 PM</p> <p>Duration 5.578s</p> <p>Compute target local</p> <p>Run ID diabetes-experiment_1586215033_2e4724ff</p> <p>Run number 2</p> <p>Script name diabetes_experiment.py</p>				<p>Label:1 3344</p> <p>Label:0 6656</p> <p>observations 10000</p>			

```
Out[8]: {'runId': 'diabetes-experiment_1586215033_2e4724ff',
  'target': 'local',
  'status': 'Finalizing',
  'startTimeUtc': '2020-04-06T23:18:56.833195Z',
  'properties': {'_azureml.ComputeTargetType': 'local',
    'ContentSnapshotId': '4b336503-4ab2-4618-b5f1-5756e4b1e48f',
    'azureml.git.repository_uri': 'https://github.com/microsoftdocs/mslearn-aml-labs',
    'mlflow.source.git.repoURL': 'https://github.com/microsoftdocs/mslearn-aml-labs',
    'azureml.git.branch': 'master',
    'mlflow.source.git.branch': 'master',
    'azureml.git.commit': 'f33990a4fa96e44ca85bedb6961c50368c92d0f7',
    'mlflow.source.git.commit': 'f33990a4fa96e44ca85bedb6961c50368c92d0f7',
    'azureml.git.dirty': 'True'},
  'inputDatasets': [],
  'runDefinition': {'script': 'diabetes_experiment.py',
    'useAbsolutePath': False,
    'arguments': [],
    'sourceDirectoryDataStore': None,
    'framework': 'Python',
    'communicator': 'None',
    'target': 'local',
    'dataReferences': {}},
  'data': {},
  'jobName': None,
  'maxRunDurationSeconds': None,
  'nodeCount': 1,
  'environment': {'name': 'Experiment diabetes-experiment Environment',
    'version': 'Autosave_2020-04-06T23:17:14Z_71f7f214',
```

- Retrieving the metrics and files it generated

```
1 [9]: # Get logged metrics
metrics = run.get_metrics()
for key in metrics.keys():
    print(key, metrics.get(key))
print('\n')
for file in run.get_file_names():
    print(file)

observations 10000
Label:0 6656
Label:1 3344

azureml-logs/60_control_log.txt
azureml-logs/70_driver_log.txt
logs/azureml/8590_azureml.log
outputs/sample.csv
```

- Retrieving an experiment by name from the workspace and iterate through its runs using the SDK

```
0: from azureml.core import Experiment, Run
diabetes_experiment = ws.experiments['diabetes-experiment']
for logged_run in diabetes_experiment.get_runs():
    print("Run ID:", logged_run.id)
    metrics = logged_run.get_metrics()
    for key in metrics.keys():
        print("-", key, metrics.get(key))

Run ID: diabetes-experiment_1586215033_2e4724ff
- observations 10000
- Label:0 6656
- Label:1 3344
Run ID: 69c0b7ca-4c21-4430-9635-c2b8ec9412b1
- observations 10000
- label distribution aml://artifactId/ExperimentRun/dcId.69c0b7ca-4c21-4430-9635-c2b8ec9412b1/label_distribution_15
86214726.png
- pregnancy categories [0, 8, 7, 9, 1, 3, 5, 2, 6, 11, 4, 13, 10, 12, 14]
- PlasmaGlucose {'stat': ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max'], 'value': [10000.0, 107.8502, 31.92090936056563, 44, 84, 105, 129, 192]}
- DiastolicBloodPressure {'stat': ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max'], 'value': [10000.0, 71.2075, 16.801478289640706, 24, 58, 72, 85, 117]}
- TricepsThickness {'stat': ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max'], 'value': [10000.0, 28.8176, 14.50648041522832, 7, 15, 31, 41, 92]}
- SerumInsulin {'stat': ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max'], 'value': [10000.0, 139.2436, 33.77791937465278, 14, 39, 85, 197, 796]}
- BMI {'stat': ['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max'], 'value': [10000.0, 31.56702174359113, 9.804365693559113, 18.20080735, 21.247426835, 31.922420785, 39.3289214475, 56.03462763]}
```

## 2. Training Models

- Connecting with the work space

```
In [1]: import azureml.core
from azureml.core import Workspace

# Load the workspace from the saved config file
ws = Workspace.from_config()
print('Ready to use Azure ML {} to work with {}'.format(azureml.core.VERSION, ws.name))

Ready to use Azure ML 1.2.0 to work with greeshma1
```

- Creating folders and files for the script

```
In [2]: import os, shutil

# Create a folder for the experiment files
training_folder = 'diabetes-training'
os.makedirs(training_folder, exist_ok=True)

# Copy the data file into the experiment folder
shutil.copy('data/diabetes.csv', os.path.join(training_folder, "diabetes.csv"))

Out[2]: 'diabetes-training/diabetes.csv'
```

- Creating training script and saving in folder

```
# calculate AUC
y_scores = model.predict_proba(X_test)
auc = roc_auc_score(y_test,y_scores[:,1])
print('AUC: ' + str(auc))
run.log('AUC', np.float(auc))

# Save the trained model in the outputs folder
os.makedirs('outputs', exist_ok=True)
joblib.dump(value=model, filename='outputs/diabetes_model.pkl')

run.complete()
```

Writing diabetes-training/diabetes\_training.py

- Using generic estimator to run the training experiment. The default environment for this estimator does not include the scikit-learn package, so this has to be added explicitly to the configuration. The conda environment is built on-demand the first time the estimator is used, and cached for future runs that use the same configuration.

### Run Properties

Status	Completed
Start Time	2020-04-06 9:51:32 PM
Duration	0:00:16
Run Id	diabetes- training_1586224290_f 8d320c3
Arguments	N/A
Accuracy	0.774
AUC	0.8484929598487486
Regularization Rate	0.01

- Getting the metrics

```
In [9]: # Get logged metrics
metrics = run.get_metrics()
for key in metrics.keys():
    print(key, metrics.get(key))
print('\n')
for file in run.get_file_names():
    print(file)

Regularization Rate 0.01
Accuracy 0.774
AUC 0.8484929598487486
```

```
azureml-logs/60_control_log.txt
azureml-logs/70_driver_log.txt
logs/azureml/8_azureml.log
outputs/diabetes_model.pkl
```

- Registering the model to track versions and fetch the model later

```
In [10]: from azureml.core import Model

# Register the model
run.register_model(model_path='outputs/diabetes_model.pkl', model_name='diabetes_model',
                    tags={'Training context':'Estimator'},
                    properties={'AUC': run.get_metrics()['AUC'], 'Accuracy': run.get_metrics()['Accuracy']})

# List registered models
for model in Model.list(ws):
    print(model.name, 'version:', model.version)
    for tag_name in model.tags:
        tag = model.tags[tag_name]
        print ('\t',tag_name, ':', tag)
    for prop_name in model.properties:
        prop = model.properties[prop_name]
        print ('\t',prop_name, ':', prop)
    print('\n')

diabetes_model version: 1
    Training context : Estimator
    AUC : 0.8484929598487486
    Accuracy : 0.774
```

- Creating folder for parameterized training script

```
In [11]: import os, shutil

# Create a folder for the experiment files
training_folder = 'diabetes-training-params'
os.makedirs(training_folder, exist_ok=True)

# Copy the data file into the experiment folder
shutil.copy('data/diabetes.csv', os.path.join(training_folder, "diabetes.csv"))

Out[11]: 'diabetes-training-params/diabetes.csv'
```

```
# calculate AUC
y_scores = model.predict_proba(X_test)
auc = roc_auc_score(y_test,y_scores[:,1])
print('AUC: ' + str(auc))
run.log('AUC', np.float(auc))

os.makedirs('outputs', exist_ok=True)
joblib.dump(value=model, filename='outputs/diabetes_model.pkl')

run.complete()

Writing diabetes-training-params/diabetes_training.py
```

- Running using framework specific estimator

run.wait\_for\_completion()

Run Properties	
Status	Completed
Start Time	2020-04-06 9:55:37 PM
Duration	0:02:39
Run Id	diabetes-training_1586224536_d74e53a6
Arguments	N/A
Accuracy	0.7736666666666666
AUC	0.8483904671874223

greeshma1 > Experiments > diabetes-training > Run 3

Run 3 ✓ Completed

↻ Refresh ▶ Resubmit ✖ Cancel

Details Metrics Images Child runs Outputs + logs Snapshot

Properties	
Status	Completed
Created	Apr 6, 2020 9:58 PM
Duration	12.43s
Compute target	local
Run ID	diabetes-training_1586224536_d74e53a6
Run number	3

- Fetching the metrics

```
In [14]: # Get logged metrics
metrics = run.get_metrics()
for key in metrics.keys():
    print(key, metrics.get(key))
print('\n')
for file in run.get_file_names():
    print(file)
```

```
Regularization Rate 0.1
Accuracy 0.7736666666666666
AUC 0.8483904671874223
```

```
azureml-logs/60_control_log.txt
azureml-logs/70_driver_log.txt
logs/azureml/8_azureml.log
outputs/diabetes_model.pkl
```

- Registering the new version of the model

```
tags={'Training context':'Parameterized SKLearn Estimator',
      'AUC': run.get_metrics()['AUC'],

# List registered models
for model in ModelClient.list(ws):
    print(model.name, 'version:', model.version)
    for tag_name in model.tags:
        tag = model.tags[tag_name]
        print ('\t',tag_name, ':', tag)
    for prop_name in model.properties:
        prop = model.properties[prop_name]
        print ('\t',prop_name, ':', prop)
    print('\n')

diabetes_model version: 2
    Training context : Parameterized SKLearn Estimator
    AUC : 0.8483904671874223
    Accuracy : 0.7736666666666666

diabetes_model version: 1
    Training context : Estimator
    AUC : 0.8484929598487486
    Accuracy : 0.774
```

### 3. Working with Data

- Connecting to the work space

```
In [1]: import azureml.core
from azureml.core import Workspace

# Load the workspace from the saved config file
ws = Workspace.from_config()
print('Ready to use Azure ML {} to work with {}'.format(azureml.__version__, ws.name))

Ready to use Azure ML 1.2.0 to work with greeshma1
```

- Determining the data stores of the workspace

```
[2]: # Get the default datastore
default_ds = ws.get_default_datastore()

# Enumerate all datastores, indicating which is the default
for ds_name in ws.datastores:
    print(ds_name, "-- Default =", ds_name == default_ds.name)

workspaceblobstore - Default = True
workspacefilestore - Default = False
```

- Uploading data into the datastore

```
        show_progress=True)

Uploading an estimated of 2 files
Uploading ./data/diabetes.csv
Uploading ./data/diabetes2.csv
Uploaded ./data/diabetes.csv, 1 files out of an estimated total of 2
Uploaded ./data/diabetes2.csv, 2 files out of an estimated total of 2
Uploaded 2 files

[3]: $AZUREML_DATAREFERENCE_468d27663e024b29b89a1391b9228327
```

- Train the model from a datastore and create folder for the same

```
[]: data_ref = default_ds.path('diabetes-data').as_download(path_on_compute='diabetes_data')
print(data_ref)

$AZUREML_DATAREFERENCE_d566947648754b6fa996dff97238918d

# Create a folder for the experiment files
experiment_folder = 'diabetes_training_from_datastore'
os.makedirs(experiment_folder, exist_ok=True)
print(experiment_folder, 'folder created.')

diabetes_training_from_datastore folder created.

# Note file saved in the outputs folder is automatically uploaded to
joblib.dump(value=model, filename='outputs/diabetes_model.pkl')

run.complete()

Writing diabetes_training_from_datastore/diabetes_training.py
```

- Setting up the script parameters to pass the file reference when we run the experiment

click to unscroll output; double click to hide	
Status	Completed
Start Time	2020-04-06 10:19:34 PM
Duration	0:00:09
Run Id	diabetes- training_1586225973_ 50bbfbc1
Arguments	N/A
Accuracy	0.7788888888888889
AUC	0.846851712258014

- Creating a dataset from the diabetes data you uploaded to the datastore, and view the first 20 records. In this case, the data is in a structured format in a CSV file, hence using tabular dataset

```
# display the first 20 rows of a pandas dataframe
tab_data_set.take(20).to_pandas_dataframe()
```

	PatientID	Pregnancies	PlasmaGlucose	DiastolicBloodPressure	TricepsThickness	SerumInsulin	BMI	DiabetesPedigree	Age	Diabetic
0	1354778	0	171	80	34	23	43.509726	1.213191	21	0
1	1147438	8	92	93	47	36	21.240576	0.158365	23	0
2	1640031	7	115	47	52	35	41.511523	0.079019	23	0
3	1883350	9	103	78	25	304	29.582192	1.282870	43	1
4	1424119	1	85	59	27	35	42.604536	0.549542	22	0
click to unscroll output; double click to hide										
6	1660149	0	133	47	19	227	21.941357	0.174160	21	0
7	1458769	0	67	87	43	36	18.277723	0.236165	26	0
8	1201647	8	80	95	33	24	26.624929	0.443947	53	1
9	1403912	1	72	31	40	42	36.889576	0.103944	26	0
10	1943830	1	88	86	11	58	43.225041	0.230285	22	0

- Creating file dataset, which creates a list of file paths in a virtual mount point, to read the data in the files.

```
In [9]: #Create a file dataset from the path on the datastore (this may take a short while
file_data_set = Dataset.File.from_files(path=(default_ds, 'diabetes-data/*.csv'))

# Get the files in the dataset
for file_path in file_data_set.to_path():
    print(file_path)

/diabetes.csv
/diabetes2.csv
```

- Registering the data stores

```
# Register the file dataset
try:
    file_data_set = file_data_set.register(workspace=ws,
                                            name='diabetes file dataset',
                                            description='diabetes files',
                                            tags = {'format':'CSV'},
                                            create_new_version=True)
except Exception as ex:
    print(ex)

print('Datasets registered')

Datasets registered
```

- Printing the list of data sets in the workspace

```
[11]: print("Datasets:")
      for dataset_name in list(ws.datasets.keys()):
          dataset = Dataset.get_by_name(ws, dataset_name)
          print("\t", dataset.name, 'version', dataset.version)
```

```
Datasets:
    diabetes file dataset version 1
    diabetes dataset version 1
```

## Train a Model from a tabular Dataset

- Create folder for the data set

```
2]: import os
```

```
# Create a folder for the experiment files
experiment_folder = 'diabetes_training_from_tab_dataset'
os.makedirs(experiment_folder, exist_ok=True)
print(experiment_folder, 'folder created')
```

```
diabetes_training_from_tab_dataset folder created
```

```
os.makedirs('outputs', exist_ok=True)
# note file saved in the outputs folder is automatically uploaded into experiment record
joblib.dump(value=model, filename='outputs/diabetes_model.pkl')

run.complete()
```

```
Writing diabetes_training_from_tab_dataset/diabetes_training.py
```

- Running the experiment

### Run Properties

Status	Completed
Start Time	2020-04-06 10:29:20 PM
Duration	0:03:18
Run Id	diabetes-training_1586226559_284debac
Arguments	N/A
Accuracy	0.7893333333333333
AUC	0.8568632924585982

## Train a Model from a File Dataset

- Creating folders for training

```
experiment_folder = 'diabetes_training_from_file_dataset'
os.makedirs(experiment_folder, exist_ok=True)
print(experiment_folder, 'folder created')

diabetes_training_from_file_dataset folder created

os.makedirs('outputs', exist_ok=True)
# note file saved in the outputs folder is automatically uploaded into experiment
joblib.dump(value=model, filename='outputs/diabetes_model.pkl')

run.complete()

Writing diabetes_training_from_file_dataset/diabetes_training.py
```

- Running the experiment

Run Properties	
Status	Completed
Start Time	2020-04-06 10:34:36 PM
Duration	0:00:12
Run Id	diabetes- training_1586226876_ 283cd6da
Arguments	N/A
Accuracy	0.7788888888888889
AUC	0.846851712258014

- Viewing the status in Azure machine learning studio

greeshma1 > Experiments > diabetes-training > Run 6

Run 6 Completed

↻ Refresh ↻ Resubmit ✖ Cancel

Details Metrics Images Child runs Outputs + logs Snapshot Raw JSON Explanations (preview)

Properties		Metrics
Status	Completed	AUC 0.846851712258014
Created	Apr 6, 2020 10:34 PM	Accuracy 0.7788888888888889
Duration	11.18s	Regularization Rate 0.1
Compute target	local	
Run ID	diabetes-training_1586226876_283cd6da	

## 4. Working with Compute

- Connecting to the work space

```

except Exception as ex:
    print(ex)
else:
    print('Dataset already registered.')

```

Dataset already registered.

- Creating folders for the logistic experiment

```

os.makedirs(experiment_folder, exist_ok=True)
print(experiment_folder, 'folder created')

diabetes_training_logistic folder created

auc = roc_auc_score(y_test,y_scores[:,1])
print('AUC: ' + str(auc))
run.log('AUC', np.float(auc))

os.makedirs('outputs', exist_ok=True)
# note file saved in the outputs folder is automatically uploaded
joblib.dump(value=model, filename='outputs/diabetes_model.pkl')

run.complete()

```

Writing diabetes\_training\_logistic/diabetes\_training.py

- Defining the environment and adding packages by using **conda** or **pip**, to ensure the experiment has access to all the libraries it requires.

```
# Create a Python environment for the experiment
diabetes_env = Environment("diabetes-experiment-env")
diabetes_env.python.user_managed_dependencies = False # Let Azure ML manage dependencies
diabetes_env.docker.enabled = True # Use a docker container

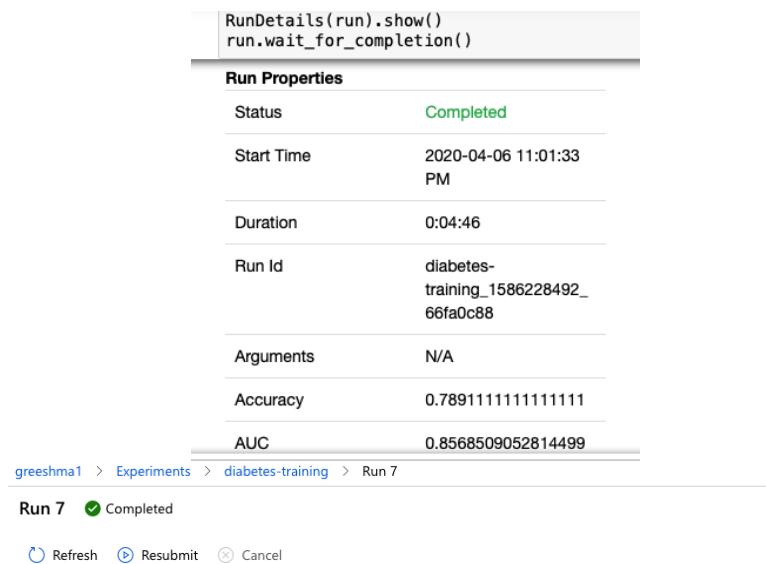
# Create a set of package dependencies (conda or pip as required)
diabetes_packages = CondaDependencies.create(conda_packages=['scikit-learn'],
                                             pip_packages=['azureml-defaults', 'azureml-dataprep[pandas]'])

# Add the dependencies to the environment
diabetes_env.python.conda_dependencies = diabetes_packages

print(diabetes_env.name, 'defined.')

diabetes-experiment-env defined.
```

- Assigning the environment created to a generic estimator, and submitting an experiment. Conda environment will also be created



The screenshot shows the Azure ML Studio interface. At the top, there is a code block with Python code for creating a environment named 'diabetes-experiment-env'. The code includes setting environment variables, creating a CondaDependencies object, and printing the environment name. Below the code, the text 'diabetes-experiment-env defined.' is displayed.

Below the code, there is a list item with the text 'Assigning the environment created to a generic estimator, and submitting an experiment. Conda environment will also be created'.

The main part of the screenshot shows the 'Run Details' for 'Run 7'. The 'Run Properties' table includes columns for Status (Completed), Start Time (2020-04-06 11:01:33 PM), Duration (0:04:46), Run Id (diabetes-training\_1586228492\_66fa0c88), Arguments (N/A), Accuracy (0.7891111111111111), and AUC (0.8568509052814499). The 'Run 7' status is 'Completed'. Below the run details, there are buttons for Refresh, Resubmit, and Cancel.

At the bottom, there is a 'Details' tab selected, showing a table with 'Properties' and 'Metrics' sections. The 'Properties' section contains fields for Status (Completed), Created (Apr 6, 2020 11:06 PM), Duration (19.59s), Compute target (local), Run ID (diabetes-training\_1586228492\_66fa0c88), and Run number (7). The 'Metrics' section contains fields for AUC (0.8568509052814499), Accuracy (0.7891111111111111), and Regularization Rate (0.1).

- Register the environment in the work space

```

: # Register the environment
diabetes_env.register(workspace=ws)

{
    "name": "diabetes-experiment-env",
    "version": "1",
    "environmentVariables": {
        "EXAMPLE_ENV_VAR": "EXAMPLE_VALUE"
    },
    "python": {
        "userManagedDependencies": false,
        "interpreterPath": "python",
        "condaDependenciesFile": null,
        "baseCondaEnvironment": null,
        "condaDependencies": {
            "channels": [
                "anaconda",
                "conda-forge"
            ],
            "dependencies": [
                "python=3.6.2",
                {
                    "pip": [

```

## Run an Experiment on a Remote Compute Target

- Utilizing compute resources to run the experiment since the local computer might not be able to handle large amount of data

```

[8]: from azureml.core.compute import ComputeTarget, AmlCompute
from azureml.core.compute_target import ComputeTargetException

cluster_name = "aml-cluster"

try:
    # Check for existing compute target
    training_cluster = ComputeTarget(workspace=ws, name=cluster_name)
    print('Found existing cluster, use it.')
except ComputeTargetException:
    # If it doesn't already exist, create it
    compute_config = AmlCompute.provisioning_configuration(vm_size='STANDARD_DS12_V2', max_nodes=4)
    training_cluster = ComputeTarget.create(ws, cluster_name, compute_config)

training_cluster.wait_for_completion(show_output=True)

Creating
Succeeded
AmlCompute wait for completion finished

Minimum number of nodes requested have been provisioned

```

- Running the experiment on the compute which was created. Since it is a free version of azure, the experiment was queued for a long time and was not scheduled to run. This step needs conda environment and assuming this takes longer time due to the free version.

```
experiment = Experiment(workspace = ws, name = 'diabetes-training')

# Run the experiment
run = experiment.submit(config=estimator)
# Show the run details while running
RunDetails(run).show()
run.wait_for_completion()
```

Run Properties	
Status	Queued
Start Time	4/8/2020 5:32:26 PM
Duration	0:05:13
Run Id	diabetes-training_1586381542_1 8da639c
Arguments	N/A

Output Logs   Auto-switch

Your job is submitted in Azure cloud and we are monitoring to get logs...

[Click here to see the run in Azure Machine Learning studio](#)

## 5. Creating an Azure Machine Learning Pipeline

- Connecting to the work space

```
n [1]: import azureml.core
from azureml.core import Workspace

# Load the workspace from the saved config file
ws = Workspace.from_config()
print('Ready to use Azure ML {} to work with {}'.format(azureml.core.VERSION, ws.name))

Ready to use Azure ML 1.2.0 to work with greeshma1
```

- Registering the data base

```

#Create a tabular dataset from the path on the datastore (this step is covered in the previous section)
tab_data_set = Dataset.Tabular.from_delimited_files(path=(def
    path))

# Register the tabular dataset
try:
    tab_data_set = tab_data_set.register(workspace=ws,
                                         name='diabetes dataset',
                                         description='diabetes data',
                                         tags = {'format':'CSV'},
                                         create_new_version=True)
    print('Dataset registered.')
except Exception as ex:
    print(ex)
else:
    print('Dataset already registered.')

```

Dataset already registered.

- Creating a folder to contain the scripts for each step

```

[3]: import os
# Create a folder for the pipeline step files
experiment_folder = 'diabetes_pipeline'
os.makedirs(experiment_folder, exist_ok=True)

print(experiment_folder)

diabetes_pipeline

auc = roc_auc_score(y_test,y_scores[:,1])
print('AUC: ' + str(auc))
run.log('AUC', np.float(auc))

# Save the trained model
os.makedirs(output_folder, exist_ok=True)
output_path = output_folder + "/model.pkl"
joblib.dump(value=model, filename=output_path)

run.complete()

```

Writing diabetes\_pipeline/train\_diabetes.py

```
# load the model
print("Loading model from " + model_folder)
model_file = model_folder + "/model.pkl"
model = joblib.load(model_file)

Model.register(workspace=run.experiment.workspace,
               model_path = model_file,
               model_name = 'diabetes_model',
               tags={'Training context':'Pipeline'})

run.complete()
```

Writing diabetes\_pipeline/register\_diabetes.py

## Prepare a Compute Environment for the Pipeline Steps

- Creating a compute target

```
compute_config = AmlCompute.provisioning_configuration(vn
                                                       m
                                                       ic
pipeline_cluster = ComputeTarget.create(ws, cluster_name,
                                         pipeline_cluster.wait_for_completion(show_output=True)

Found existing cluster, use it.
Succeeded
AmlCompute wait for completion finished

Minimum number of nodes requested have been provisioned
```

- Compute will require a Python environment with the necessary package dependencies installed

```

# Add the dependencies to the environment
diabetes_env.python.conda_dependencies = diabetes_packages

# Register the environment (just in case you want to use it again)
diabetes_env.register(workspace=ws)
registered_env = Environment.get(ws, 'diabetes-pipeline-env')

# Create a new runconfig object for the pipeline
pipeline_run_config = RunConfiguration()

# Use the compute you created above.
pipeline_run_config.target = pipeline_cluster

# Assign the environment to the run configuration
pipeline_run_config.environment = registered_env

print ("Run configuration created.")

```

Run configuration created.

## Create and Run a Pipeline

- Defining pipeline steps

```

# Create a PipelineData (temporary Data Reference) for the model folder
model_folder = PipelineData("model_folder", datastore=ws.get_default_datastore())

estimator = Estimator(source_directory=experiment_folder,
                      compute_target = pipeline_cluster,
                      environment_definition=pipeline_run_config.environment,
                      entry_script='train_diabetes.py')

# Step 1, run the estimator to train the model
train_step = EstimatorStep(name = "Train Model",
                           estimator=estimator,
                           estimator_entry_script_arguments=['--output_folder', model_folder],
                           inputs=[diabetes_ds.as_named_input('diabetes_train')],
                           outputs=[model_folder],
                           compute_target = pipeline_cluster,
                           allow_reuse = True)

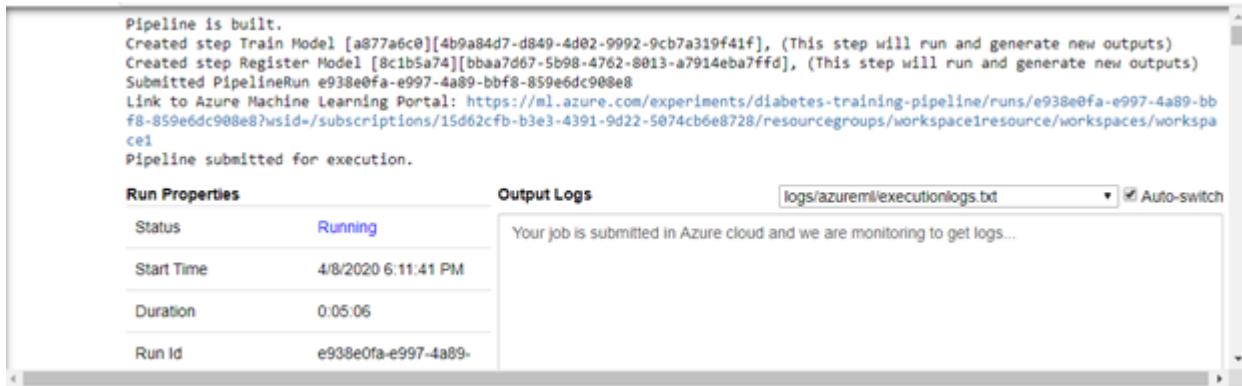
# Step 2, run the model registration script
register_step = PythonScriptStep(name = "Register Model",
                                 source_directory = experiment_folder,
                                 script_name = "register_diabetes.py",
                                 arguments = ['--model_folder', model_folder],
                                 inputs=[model_folder],
                                 compute_target = pipeline_cluster,
                                 runconfig = pipeline_run_config,
                                 allow_reuse = True)

print("Pipeline steps defined")

```

Pipeline steps defined

- Running the pipeline. I left this to run for almost two days but I believe this will not be completed since it is a free azure account. I am assuming these jobs would be set priority and since it is a free account, it was not fully finished.



## 6. Creating a Real-Time Inferencing Service

- Connecting to work space

```
In [2]: import azureml.core
from azureml.core import Workspace

# Load the workspace from the saved config file
ws = Workspace.from_config()
print('Ready to use Azure ML {} to work with {}'.format(azureml.core.VERSION, ws.name))

Ready to use Azure ML 1.2.0 to work with greeshma1
```

- Training the model to deploy

```
auc = roc_auc_score(y_test, y_scores, average='macro')
print('AUC: ' + str(auc))
run.log('AUC', np.float(auc))

# Save the trained model
model_file = 'diabetes_model.pkl'
joblib.dump(value=model, filename=model_file)
run.upload_file(name = 'outputs/' + model_file, path_or_stream = './' + model_file)

# Complete the run
run.complete()

# Register the model
run.register_model(model_path='outputs/diabetes_model.pkl', model_name='diabetes_model',
                   tags={'Training context':'Inline Training'},
                   properties={'AUC': run.get_metrics()['AUC'], 'Accuracy': run.get_metrics()['Accuracy']})

print('Model trained and registered.')

Starting experiment: diabetes-training
Loading Data...
Training a decision tree model
Accuracy: 0.8936666666666667
AUC: 0.8808423304642021
Model trained and registered.
```

## Deploy a Model as a Web Service

- Determining what models you have registered in the workspace

```
for model in Model.list(ws):
    print(model.name, 'version:', model.version)
    for tag_name in model.tags:
        tag = model.tags[tag_name]
        print ('\t',tag_name, ':', tag)
    for prop_name in model.properties:
        prop = model.properties[prop_name]
        print ('\t',prop_name, ':', prop)
    print('\n')

diabetes_model version: 3
    Training context : Inline Training
    AUC : 0.8808423304642021
    Accuracy : 0.8936666666666666

diabetes_model version: 2
    Training context : Parameterized SKLearn Estimator
    AUC : 0.8483904671874223
    Accuracy : 0.7736666666666666

diabetes_model version: 1
    Training context : Estimator
    AUC : 0.8484929598487486
    Accuracy : 0.774
```

- Getting the model which has to be deployed

```
: model = ws.models['diabetes_model']
print(model.name, 'version', model.version)

diabetes_model version 3
```

- Creating folder for code and configuration files

```
[6]: import os

folder_name = 'diabetes_service'

# Create a folder for the web service files
experiment_folder = './' + folder_name
os.makedirs(folder_name, exist_ok=True)

print(folder_name, 'folder created.')

diabetes_service folder created.
```

- The web service to deploy the model will need some Python code to load the input data, get the model from the workspace, and generate and return predictions. Saving this code in an entry script that will be deployed to the web service

```
model = joblib.load(model_path)

# Called when a request is received
def run(raw_data):
    # Get the input data as a numpy array
    data = np.array(json.loads(raw_data)['data'])
    # Get a prediction from the model
    predictions = model.predict(data)
    # Get the corresponding classname for each prediction (0 or 1)
    classnames = ['not-diabetic', 'diabetic']
    predicted_classes = []
    for prediction in predictions:
        predicted_classes.append(classnames[prediction])
    # Return the predictions as JSON
    return json.dumps(predicted_classes)
```

```
Writing diabetes_service/score_diabetes.py
```

- Web service will be hosted in a container, and the container will need to install any required Python dependencies when it gets initialized. In this case, scoring code requires scikit-learn, so creating a .yml file that tells the container host to install this into the environment

```
env_file = folder_name + "/diabetes_env.yml"
with open(env_file,"w") as f:
    f.write(myenv.serialize_to_string())
print("Saved dependency info in", env_file)

# Print the .yml file
with open(env_file,"r") as f:
    print(f.read())
```

```
Saved dependency info in diabetes_service/diabetes_env.yml
# Conda environment specification. The dependencies defined in this file will
# be automatically provisioned for runs with userManagedDependencies=False.

# Details about the Conda environment file format:
# https://conda.io/docs/user-guide/tasks/manage-environments.html#create-env-file-manually

name: project_environment
dependencies:
    # The python interpreter version.
    # Currently Azure ML only supports 3.5.2 and later.
    - python=3.6.2

    - pip:
        # Required packages for AzureML execution, history, and data preparation.
        - azureml-defaults

    - scikit-learn
channels:
```

- Deploying the model

```
]: from azureml.core.webservice import AciWebservice
from azureml.core.model import InferenceConfig

# Configure the scoring environment
inference_config = InferenceConfig(runtime="python",
                                    source_directory=folder_name,
                                    entry_script="score_diabetes.py",
                                    conda_file="diabetes_env.yml")

deployment_config = AciWebservice.deploy_configuration(cpu_cores=1, memory_gb=1)

service_name = "diabetes-service"

service = Model.deploy(ws, service_name, [model], inference_config, deployment_config)

service.wait_for_deployment(True)
print(service.state)
```

Running.....  
 Succeeded  
 ACI service creation operation finished, operation "Succeeded"  
 Healthy

```
print(service.state)
print(service.get_logs())

# If you need to make a change and redeploy, you may need to delete
#service.delete()
```

Healthy

```
2020-04-07T04:17:11,760622098+00:00 - rsyslog/run
2020-04-07T04:17:11,766928615+00:00 - gunicorn/run
2020-04-07T04:17:11,771446726+00:00 - nginx/run
/usr/sbin/nginx: /azureml-envs/azureml_4b824bcb98517d791c41923f24d6
ion available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_4b824bcb98517d791c41923f24d6
ion available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_4b824bcb98517d791c41923f24d6
available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_4b824bcb98517d791c41923f24d6
available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_4b824bcb98517d791c41923f24d6
available (required by /usr/sbin/nginx)
2020-04-07T04:17:11,781348352+00:00 - iot-server/run
EdgeHubConnectionString and IOTEDGE_IOTHUBHOSTNAME are not set. Exi
2020-04-07T04:17:12,012840151+00:00 - iot-server/finish 1 0
2020-04-07T04:17:12,012840151+00:00 - iot-server/finish 1 0
```

```
: for webservice_name in ws.webservices:
    print(webservice_name)
```

diabetes-service

- The service deployed can be consumed it from a client application

```

import json

x_new = [[2,180,74,24,21,23.9091702,1.488172308,22]]
print ('Patient: {}'.format(x_new[0]))

# Convert the array to a serializable list in a JSON document
input_json = json.dumps({"data": x_new})

# Call the web service, passing the input data (the web service
predictions = service.run(input_data = input_json)

# Get the predicted class - it'll be the first (and only) one.
predicted_classes = json.loads(predictions)
print(predicted_classes[0])

```

```

Patient: [2, 180, 74, 24, 21, 23.9091702, 1.488172308, 22]
diabetic

```

- Sending multiple patient observations to the service, and get back a prediction for each one

```

: import json

# This time our input is an array of two feature arrays
x_new = [[2,180,74,24,21,23.9091702,1.488172308,22],
          [0,148,58,11,179,39.19207553,0.160829008,45]]

# Convert the array or arrays to a serializable list in a JSON document
input_json = json.dumps({"data": x_new})

# Call the web service, passing the input data
predictions = service.run(input_data = input_json)

# Get the predicted classes.
predicted_classes = json.loads(predictions)

for i in range(len(x_new)):
    print ("Patient {}".format(x_new[i]), predicted_classes[i] )

```

```

Patient [2, 180, 74, 24, 21, 23.9091702, 1.488172308, 22] diabetic
Patient [0, 148, 58, 11, 179, 39.19207553, 0.160829008, 45] not-diabetic

```

- Determining the URL to which these applications must submit their requests

```
endpoint = service.scoring_uri
print(endpoint)

http://15f8aded-67fc-4c6a-9ce9-917ade620918.eastus2.azurecontainer.io/score
```

- An application can simply make an HTTP request, sending the patient data in JSON (or binary) format, and receive back the predicted class(es)

```
import requests
import json

x_new = [[2,180,74,24,21,23.9091702,1.488172308,22],
          [0,148,58,11,179,39.19207553,0.160829008,45]]

# Convert the array to a serializable list in a JSON document
input_json = json.dumps({"data": x_new})

# Set the content type
headers = { 'Content-Type':'application/json' }

predictions = requests.post(endpoint, input_json, headers = headers)
predicted_classes = json.loads(predictions.json())

for i in range(len(x_new)):
    print ("Patient {}".format(x_new[i]), predicted_classes[i] )

Patient [2, 180, 74, 24, 21, 23.9091702, 1.488172308, 22] diabetic
Patient [0, 148, 58, 11, 179, 39.19207553, 0.160829008, 45] not-diabetic
```

## 7. Using Automated Machine Learning

- Connecting to work space

```
[1]: import azureml.core
from azureml.core import Workspace

# Load the workspace from the saved config file
ws = Workspace.from_config()
print('Ready to use Azure ML {} to work with {}'.format(azureml.core.VERSION, ws.name))

Ready to use Azure ML 1.2.0 to work with greeshma1
```

- Running configuration that includes the required packages for the experiment environment, configuration settings that specifies how many combinations to try, which metric to use when evaluating models

```
[2]: import pandas as pd
from azureml.train.automl import AutoMLConfig

# Load the data
train_data = pd.read_csv('data/diabetes.csv')

automl_config = AutoMLConfig(name='Automated ML Experiment',
                             task='classification',
                             compute_target='local',
                             training_data = train_data,
                             n_cross_validations = 2,
                             label_column_name='Diabetic',
                             iterations=6,
                             primary_metric = 'AUC_weighted',
                             max_concurrent_iterations=3,
                             featurization='auto'
                             )

print("Ready for Auto ML run.")
```

Ready for Auto ML run.

- Running the automated machine learning experiment

```
In [3]: from azureml.core.experiment import Experiment
from azureml.widgets import RunDetails

print('Submitting Auto ML experiment...')
automl_experiment = Experiment(ws, 'diabetes_automl')
automl_run = automl_experiment.submit(automl_config)
RunDetails(automl_run).show()
automl_run.wait_for_completion()
```

Submitting Auto ML experiment...

AutoML\_70f46d1f-1b5a-495a-b772-f5738db367ce:

Status: Completed

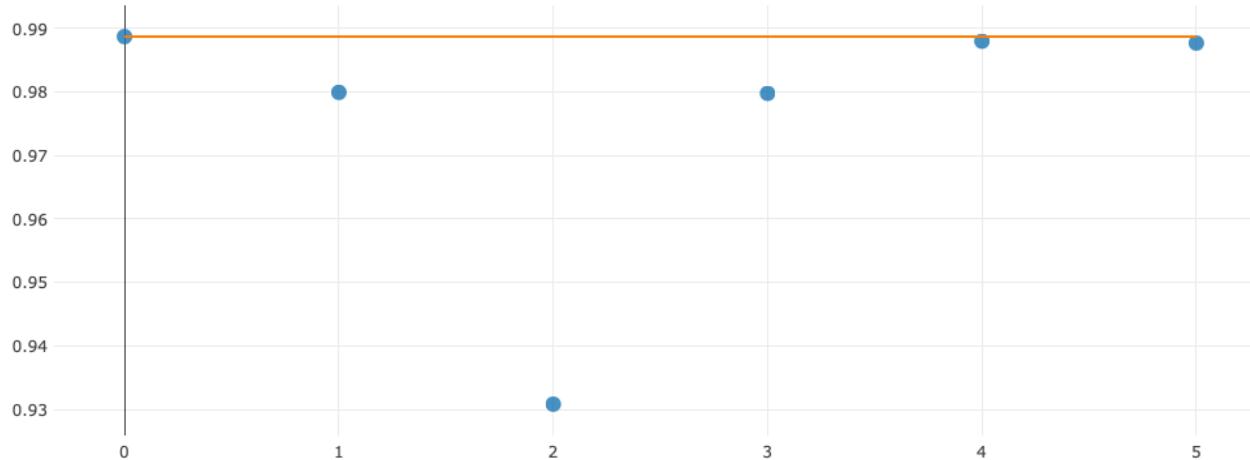


click to scroll output; double click to hide

1 2 3 4 5

Iteration	Pipeline	Iteration metric	Best metric	Status	Duration	Started
0	MaxAbsScaler, LightGBM	0.98873079	0.98873079	Completed	0:00:22	Apr 7, 2020 12:37 AM
4	VotingEnsemble	0.988801063	0.98873079	Completed	0:00:13	Apr 7, 2020 12:38 AM
5	StackEnsemble	0.9877102	0.98873079	Completed	0:00:12	Apr 7, 2020 12:39 AM
1	MaxAbsScaler, LightGBM	0.97994864	0.98873079	Completed	0:00:21	Apr 7, 2020 12:37 AM

### AutoML Run with metric : AUC\_weighted



greeshma1 > Experiments > diabetes\_automl > Run 1

Run 1 Completed Switch to old experience

⟳ Refresh ✖ Cancel

Details Data guardrails Models Logs Outputs

**Model summary**

Algorithm name  
MaxAbsScaler, LightGBM

AUC weighted  
0.9887307944471648 [View all other metrics](#)

Registered models  
[diabetes\\_model\\_automl:1](#)

Deploy status  
No deployment yet

Deploy best model View model details

Download best model

**Run summary**

Task type  
Classification [View all run settings](#)

Primary metric  
AUC weighted

Run status  
Completed

Experiment name  
diabetes\_automl

Run ID  
AutoML\_70f46d1f-1b5a-495a-b772-f5738db367ce

Input datasets  
--

- Determining the best model

```
4]: best_run, fitted_model = automl_run.get_output()
print(best_run)
print(fitted_model)
best_run_metrics = best_run.get_metrics()
for metric_name in best_run_metrics:
    metric = best_run_metrics[metric_name]
    print(metric_name, metric)

Run(Experiment: diabetes_automl,
Id: AutoML_70f46d1f-1b5a-495a-b772-f5738db367ce_0,
Type: None,
Status: Completed)
Pipeline(memory=None,
    steps=[('datatransformer', DataTransformer(enable_dnn=None, enable_feature_sweeping=None,
                                                feature_sweeping_config=None, feature_sweeping_timeout=None,
                                                featurization_config=None, force_text_dnn=None,
                                                is_cross_validation=None, is_onnx_compatible=None, logger=None,
                                                obser...    silent=True, subsample=1.0, subsample_for_bin=200000,
                                                subsample_freq=0, verbose=-10))])
norm_macro_recall 0.8794903988906442
average_precision_score_weighted 0.9893133642855426
average_precision_score_macro 0.9868730917793174
precision_score_macro 0.9432823058330543
```

- Viewing the steps in a model

```
: for step in fitted_model.named_steps:
    print(step)

datatransformer
MaxAbsScaler
LightGBMClassifier
```

- Registering the best performing model

```

for tag_name in model.tags:
    tag = model.tags[tag_name]
    print ('\t',tag_name, ':', tag)
for prop_name in model.properties:
    prop = model.properties[prop_name]
    print ('\t',prop_name, ':', prop)
print('\n')

diabetes_model_automl version: 1
    Training context : Auto ML
    AUC : 0.9887307944471648
    Accuracy : 0.9480999999999999

diabetes_model version: 3
    Training context : Inline Training
    AUC : 0.8808423304642021
    Accuracy : 0.8936666666666667

diabetes_model version: 2
    Training context : Parameterized SKLearn Estimator
    AUC : 0.8483904671874223
    Accuracy : 0.7736666666666666

diabetes_model version: 1
    Training context : Estimator
    AUC : 0.8484929598487486
    Accuracy : 0.774

```

- Many experiments were run as a part of implementing machine learning models in different ways. Out of all the different ways, Auto ML was giving the highest accuracy of 94% for the given image data set.

Experiment	Latest run	Last submitted	Created	Created by	Run types
diabetes-training-pipeline	3	Apr 7, 2020 9:54 ...	Apr 6, 2020 11:3...	G. Gopal	Pipeline
diabetes-training	13	Apr 7, 2020 9:52 ...	Apr 6, 2020 8:34 ...	G. Gopal	Script
diabetes_automl	1	Apr 7, 2020 12:3...	Apr 7, 2020 12:3...	G. Gopal	Automated ML
diabetes-experiment	2	Apr 6, 2020 7:17 ...	Apr 6, 2020 7:11 ...	G. Gopal	Script